





Algorithms and Data Structures

Week 10

Shortest Path



Vertex	Visited	Distance	Previous
A	F	∞	\emptyset
B	F	∞	\emptyset
C	F	∞	\emptyset
D	F	∞	\emptyset
E	F	14	H
F	F	∞	\emptyset
G	F	19	H
H	T	8	K
I	F	10	H
J	F	17	K
K	T	0	\emptyset
L	F	16	K



Acknowledgement:

The slide is modified from

- Fan Rui's CS140@ShanghaiTech, Fall 2018
- TA discussion CS101 @ShanghaiTech, Fall 2018
 - Lecture CS101 @ShanghaiTech, Fall 2019
- KeWei Tu's CS181@ShanghaiTech, Fall 2018

Shortest Path

Today's topic:

- SSSP: The Dijkstra's algorithm
 - Introduce the motivation of relaxation and what it is
 - Introduce the algorithm and its time complexity with binary heap
 - Introduce the correctness of it(prepare for Prof. Dengji, Zhao 's algorithm section)
- APSP: The Floyd-Warshall algorithm
 - Introduce the algorithm
- A*: An improvement of Dijkstra's algorithm
 - A quick proof
 - Applications

Shortest Path

Today's topic:


- SSSP: The Dijkstra's algorithm
 - Introduce the motivation of relaxation and what it is
 - Introduce the algorithm and its time complexity with binary heap
 - Introduce the correctness of it(prepare for Prof. Dengji, Zhao 's algorithm section)
- APSP: The Floyd-Warshall algorithm
 - Introduce the algorithm
- A*: An improvement of Dijkstra's algorithm
 - A quick proof
 - Applications

Shortest Path

- Given a weighted directed graph, one common problem is finding the shortest path between two given vertices.
- This is different from the minimum spanning tree.

SSSP: Dijkstra's algorithm

Idea *shortest distance from s to $v \leq \text{some distance from } s \text{ to } u + \text{distance from } u \text{ to } v$*



Given source s and node v

Let $\delta(s, v)$ be length of a shortest path from s to v

Let $d(s, v)$ be an estimate of $\delta(s, v)$, where $d(s, v) \geq \delta(s, v)$.

$d(s, v)$ is the shortest distance as for our best knowledge!

Intuitively, we can update $d(s, v)$ by iterations to “push” it to $\delta(s, v)$.


How? By **relaxation**!

In lecture, Prof. Yuyao, Zhang has guided you how it works by specific example of Dijkstra's algorithm(find the shortcut!)

SSSP: Dijkstra's algorithm

Relaxation Given a neighbor u of v ,

$$d(s, v) \leftarrow \min(d(s, v), d(s, u) + w(u, v))$$



RELAX(u, v, w)

- 1 **if** $v.d > u.d + w(u, v)$
- 2 $v.d = u.d + w(u, v)$
- 3 $v.\pi = u$

- $\delta(s, v)$ be length of a shortest path from s to v
- $d(s, v)$ is the shortest distance we known

Relaxation compares two paths and picks a better one:

- A shortest path we known(an estimate)
- Another one path which passes u , then reaches v directly

SSSP: Dijkstra's algorithm

SSSP: single source shortest paths

We will iterate $|V|$ times:

- Find the unvisited vertex v that has a minimum distance to it
- Mark it as visited
- Consider its every adjacent vertex w that is unvisited:
 - Is the distance to v plus the weight of the edge (v, w) less than our currently known shortest distance to w ?
 - If so, update the shortest distance(**relaxation**) to w and record v as the previous pointer

Continue iterating until all vertices are visited or **all remaining vertices have a distance of infinity**

DIJKSTRA(G, w, s)

```
1 INITIALIZE-SINGLE-SOURCE( $G, s$ )
2  $S = \emptyset$ 
3  $Q = G.V$ 
4 while  $Q \neq \emptyset$ 
5      $u = \text{EXTRACT-MIN}(Q)$ 
6      $S = S \cup \{u\}$ 
7     for each vertex  $v \in G.Adj[u]$ 
8         RELAX( $u, v, w$ )
```


SSSP: Dijkstra's algorithm

Questions:

- What if at some point, all unvisited vertices have a distance ∞ ?
 - This means that the graph is unconnected
 - We have found the shortest paths to all vertices in the connected subgraph containing the source vertex
- What if we just want to find the shortest path between vertices v_j and v_k ?
 - Apply the same algorithm, but stop when we are visiting vertex v_k
- Does the algorithm change if we have a directed graph?
 - No


SSSP: Dijkstra's algorithm

Questions:

- How to re-construct the shortest path?
 - Recall we record previous pointer when each relaxation happens

Shortest path tree There exists a tree T rooted at s such that the shortest path from s to any node v lies in T .

- Each node v has a parent in the tree.
- By following parent pointers starting from v , we find shortest path from s to v .



SSSP: Dijkstra's algorithm

The initialization requires $\Theta(|V|)$ memory and run time

- Iterating through the table requires $\Theta(|V|)$ time
- Each time we find a vertex, we must check all of its neighbors
 - With an adjacency matrix, the run time is
 - $\Theta(|V|(|V|(\text{find closest}) + |V|(\text{relaxation}))) = \Theta(|V|^2)$
 - With an adjacency list, the run time is
 - $\Theta(|V|^2(\text{find closest}) + |E|(\text{relaxation})) = \Theta(|V|^2)$ as $|E| = O(|V|^2)$

SSSP: Dijkstra's algorithm

Can we do better?

- How about using a priority queue to find the closest vertex?
 - Assume we are using a binary heap
 - Thus, the total run time is $O(|V| \ln(|V|) + |E| \ln(|V|))$
 - $O(|V| \ln(|V|))$: find closest vertex and pop it (recall time complexity for a binary heap)
 - $O(|E| \ln(|V|))$: for each edge, we do a relaxation, which influences the estimated distance for a vertex. We should maintain the heap.
 - Naïve solution: delete it, decrease it then insert it.
 - It is called **decrease the key** in heap context (the key is the estimated distance in shortest path context. Relaxation decreases the estimation)

SSSP: Dijkstra's algorithm

Can we do better then better?

Dijkstra's algorithm for positive edge weights, in $O((|V| + |E|) \log |V|)$ time by binary heap(has discussed in last page).


Improve to $O(|V| \log |V| + |E|)$ using **Fibonacci heap**

Using a Fibonacci heap, the $O(|E|)$ decrease keys each take $O(1)$, on average, which is used to be $\log |V|$ (omit details here. Clarifying it refers to an advanced topic: amortization analysis).

SSSP: Dijkstra's algorithm

Negative cycles

- A **negative weight cycle** is a cycle in the graph, s.t. the sum of all weights on cycle is negative.
- If a graph has a negative weight cycle reachable from the source, then shortest paths are not well defined
 - We can repeatedly go around the cycle to get arbitrarily short paths
- Dijkstra's algorithm assumes all weights are nonnegative(why?)
- Counter example:




SSSP: Dijkstra's algorithm

Correctness

Idea Order nodes by **increasing distance** from source, as s, v_1, \dots, v_n . s is source. The larger index indicates larger distance of that vertex to the source.

$$\delta(v_1, s) \leq \delta(v_2, s) \leq \dots \leq \delta(v_n, s)$$

Lemma In round k of Dijkstra's algorithm, node v_k is settled (found shortest path, visited), and $S = \{s, v_1, \dots, v_k\}$. S is the set of visited vertexes.



SSSP: Dijkstra's algorithm

Correctness

Lemma In round k of Dijkstra's algorithm, node v_k is settled (founded shortest path, visited), and $S = \{s, v_1, \dots, v_k\}$. S is the set of visited vertexes.

Proof

1. In round 0, s is settled.
2. Assume lemma on round $k-1$ establishes, then we prove lemma on k (induction).

Consider shortest path from s to v_k , and let u be node preceding v_k in path.


- Then $\delta(s, v_k) = \delta(s, u) + w(u, v_k)$

Since $w(.) > 0$, the equation indicates that:

- $\delta(s, v_k) \geq \delta(s, u)$

By increasing distance notation, $\exists j < k: u = v_j$.

Otherwise, vertex k is not connected to source (trivial, not consider)



SSSP: Dijkstra's algorithm

Correctness


Lemma In round k of Dijkstra's algorithm, node v_k is settled (found shortest path, visited), and $S = \{s, v_1, \dots, v_k\}$. S is the set of visited vertexes.

Proof (cont.)

2. Assume lemma on round $k-1$ establishes, then we prove lemma on k (induction).

By induction, $v_j.d$ (the estimated distance) = $\delta(s, v_j)$ (the optimal distance).

- v_j is processed in round j (induction, $j < k$). v_k is its neighbor and being relax:
 - $v_k.d \leq v_j.d (= \delta(s, v_j)) + w(v_j, v_k)$ since $v_k.d = \min(v_k.d, v_k.d + w(v_j, v_k))$.
- Since $u = v_j$ therefore, $v_k.d \leq \delta(s, u) + w(v_k, u) = \delta(s, v_k) \Rightarrow v_k.d = \delta(s, v_k)$ after relaxation in round k .



SSSP: Dijkstra's algorithm

Correctness

Lemma In round k of Dijkstra's algorithm, node v_k is settled (found shortest path, visited), and $S = \{s, v_1, \dots, v_k\}$. S is the set of visited vertexes.

Proof (cont.)


2. Assume lemma on round $k-1$ establishes, then we prove lemma on k (induction).

$v_k.d = \delta(s, v_k)$ after relaxation in round k .

And in round k , v_k is sure to be selected.

- Since by induction, $V - S = \{v_k, v_{k+1}, \dots\}$ in round k before selection of closet vertex. $S = \{s, v_1, \dots, v_k\}$

q.e.d



Recursively apply Lemma, then correctness of Dijkstra's algorithm is guaranteed

Shortest Path

Today's topic:

- SSSP: The Dijkstra's algorithm
 - Introduce the motivation of relaxation and what it is
 - Introduce the algorithm and its time complexity with binary heap
 - Introduce the correctness of it(prepare for Prof. Dengji, Zhao 's algorithm section)
- APSP: The Floyd-Warshall algorithm
 - Introduce the algorithm
- A*: An improvement of Dijkstra's algorithm
 - A quick proof
 - Applications

APSP: Floyd-Warshall algorithm

ASSP: all pairs shortest paths

First, let's consider only edges that connect vertices directly:

$$d_{i,j}^{(0)} = \begin{cases} 0 & \text{If } i = j \\ w_{i,j} & \text{If there is an edge from } i \text{ to } j \\ \infty & \text{Otherwise} \end{cases}$$

Here, $w_{i,j}$ is the weight of the edge connecting vertices i and j

- Note, this can be a directed graph; *i.e.*, it may be that

APSP: Floyd-Warshall algorithm

The calculation is straight forward:

```
for ( int i = 0; i < num_vertices; ++i ) {  
    for ( int j = 0; j < num_vertices; ++j ) {  
        d[i][j] = std::min( d[i][j], d[i][k-1] + d[k-1][j] );  
    }  
}
```

APSP: Floyd-Warshall algorithm

What Is the Shortest Path?

Let us store the next vertex in the shortest path. Initially:

$$p_{i,j} = \begin{cases} \emptyset & \text{If } i = j \\ j & \text{If there is an edge from } i \text{ to } j \\ \emptyset & \text{Otherwise} \end{cases}$$

APSP: Floyd-Warshall algorithm

When we find a shorter path, update the next node(relaxation):

$$p_{i,j} = p_{i,k}$$

```
// Initialize the matrix p
// ...

for ( int k = 0; k < num_vertices; ++k ) {
    for ( int i = 0; i < num_vertices; ++i ) {
        for ( int j = 0; j < num_vertices; ++j ) {
            if ( d[i][j] > d[i][k] + d[k][j] ) {
                p[i][j] = p[i][k];
                d[i][j] = d[i][k] + d[k][j];
            }
        }
    }
}
```

Takes $O(n^3)$ time overall.

Shortest Path

Today's topic:

- SSSP: The Dijkstra's algorithm
 - Introduce the motivation of relaxation and what it is
 - Introduce the algorithm and its time complexity with binary heap
 - Introduce the correctness of it(prepare for Prof. Dengji, Zhao 's algorithm section)
- APSP: The Floyd-Warshall algorithm
 - Introduce the algorithm
- A*: An improvement of Dijkstra's algorithm
 - A quick proof
 - Applications

A* Search

SSSP is to find shortest paths from source to all vertexes in graph. When there is a **goal** vertex and regard weights as **cost**, it can be reframed as a searching problem:

Search the goal with lowest cost

(Dijkstra halts when visit the goal vertex)





A* Search

The strategy employed by Dijkstra: expand the cheapest node first is also called:

Uniform Cost Search

Strategy: expand a
cheapest node first:

Fringe is a priority queue
(priority: cumulative cost)





In some context, UCS can access visited vertexes(trivial). 'Visit' is called 'expand' here.

A* Search

Uniform Cost Search




- The good: UCS is complete and optimal!
- The bad:
 - Explores options in every “direction”
 - No information about goal location
- We'll fix that soon!



A* Search

Search heuristics

- A heuristic is:
 - A function that *estimates* how close a state is to a goal
 - Designed for a particular search problem
 - Examples: Manhattan distance, Euclidean distance for pathing




A* Search


Greedy Search

Strategy: visit(expand) the node that you think is closest to a goal state (lowest heuristics value).

The ideal scenario: best-first takes you straight to the goal.



Worst-case: like a badly-guided DFS



A* Search

A* Search



UCS



Greedy

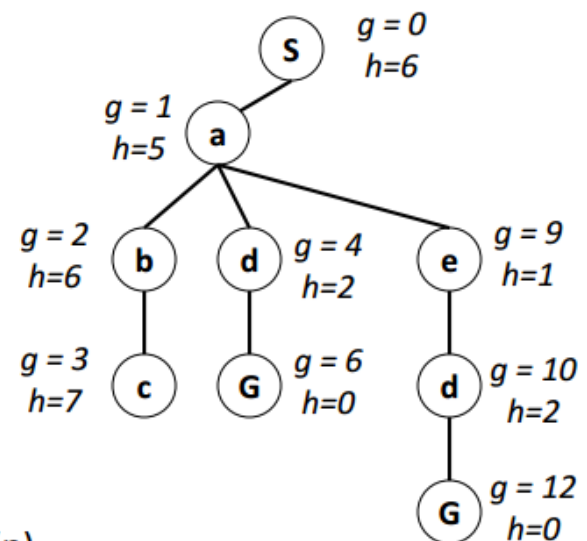



A*

A* Search

UCS v.s. Greedy

- Uniform-cost orders by path cost, or *backward cost* $g(n)$
- Greedy orders by goal proximity, or *forward cost* $h(n)$




- A* Search orders by the sum: $f(n) = g(n) + h(n)$

Example: Teg Grenager

A* Search

Optimal?



- What went wrong?
- Over-estimated goal cost

A* Search



Admissible Heuristics

- A heuristic h is *admissible* (optimistic) if:

$$0 \leq h(n) \leq h^*(n)$$

where $h^*(n)$ is the true cost to a nearest goal

- Examples:



- Coming up with admissible heuristics is most of what's involved in using A* in practice.

A* Search


Optimality of A* with admissible heuristics

Assume:

- A is an optimal goal node
- B is a suboptimal goal node
- h is admissible

Claim:

- A will exit the fringe before B



Remark1: fringe is the vertexes(nodes) adjacent to visited vertexes


Remark2: A, B here may mean different paths to goal(relaxed from different parents). Also, goal could also be a set of nodes in A* setting.

A* Search

Optimality of A* with admissible heuristics

Proof:

- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B
 - $f(n)$ is less or equal to $f(A)$



$$f(n) = g(n) + h(n)$$

$$f(n) \leq g(A)$$

$$g(A) = f(A)$$

Definition of f-cost

Admissibility of h


$h = 0$ at a goal

A* Search

Optimality of A* with admissible heuristics

Proof:

- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B
 - $f(n)$ is less or equal to $f(A)$
 - $f(A)$ is less than $f(B)$



$$g(A) < g(B)$$

$$f(A) < f(B)$$

B is suboptimal


$h = 0$ at a goal

A* Search

Optimality of A* with admissible heuristics

Proof:

- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B
 1. $f(n)$ is less or equal to $f(A)$
 2. $f(A)$ is less than $f(B)$
 3. n expands before B
- All ancestors of A expand before B
- A expands before B
- A* search is optimal




$$f(n) \leq f(A) < f(B)$$


A* Search

Optimality of A* with admissible heuristics

- Uniform-cost expands equally in all “directions”



- A* expands mainly toward the goal, but does hedge its bets to ensure optimality



A* Search

Applications


- Video games
- Pathing / routing problems
- Resource planning problems
- Robot motion planning
- Language analysis
- Machine translation
- Speech recognition
- ...




A* Search

Applications


- Heuristic: Number of tiles misplaced
- $h(\text{start}) = 8$
- Is it admissible?
- This is a *relaxed-problem* heuristic



Start State



Goal State



Average nodes expanded when the optimal path has...			
	...4 steps	...8 steps	...12 steps
UCS	112	6,300	3.6×10^6
TILES	13	39	227

Thank you !

Goodbye for data structure section and say hello
world to algorithm part!